SETLIST

START HERE · NOW · HUM QUALITY · AGENT AND RELIABILITY · VALIDATION CORRECTNESS · RETRIEVAL AND DATA · TESTS AND HOUSEKEEPING · LATER · IDEAS

NEXT STEPS

StrudelAssistant Setlist

What comes next for the web app, now that the iOS shell is gone. Three steps need you in person and nothing that calls the API can be tested until the first one is done. The rest can be handed off in almost any order.

follows **feat/hum-to-pattern** at **9814eb9** **27** steps · **6** ideas · **6** need you 14 October 2026

START HERE

Add credits, then hum into it

The OpenAI account is out of credits, so the model cannot be called. Once it is topped up, a real voice is the one test that can overturn the tuning, and it takes fifteen minutes.

1. Open the page in Safari or Chrome. Embedded previews block the microphone.
2. Press the red dot, hum, press it again.
3. Judge the panel labelled *As hummed* first. It involves no model.

```
uv run python app/server.py
# then open http://localhost:8000
```

For each take, note three things: did held notes shatter into repeats, did repeated notes merge into one, and was the tempo right, half or double.

NOW 3 steps	**HUM QUALITY** 4 steps	**AGENT AND RELIABILITY** 5 steps	**VALIDATION CORRECTNESS** 3 steps
RETRIEVAL AND DATA 4 steps	**TESTS AND HOUSEKEEPING** 5 steps	**LATER** 3 steps	**IDEAS** 6, not commitments

NOW

Blocked on you

Credits and a real voice come first: nothing that calls the API can be tested without credits, and every accuracy number so far comes from synthesized audio.

01 YOU 5 min

Add OpenAI credits

The account returned `insufficient_quota` on 14 October, so chat, hum and the benchmark all fail until it is topped up.

DONE WHEN	A one-line tool-calling request to the chat model succeeds.
UNBLOCKS	02, 03 and everything that needs the API

02 YOU small

Try `gpt-5.6-luna` in place of `gpt-5.4-mini`

Third-party trackers put Luna near $0.20 / $1.20 per million tokens against $0.75 / $4.50 for the current model. That is unconfirmed, and so is whether Luna works on the Chat Completions path.

HOW	Check tool calling with `init_chat_model`, then rerun the five-run melody benchmark recorded in the comment in `core.py`.
DONE WHEN	Luna scores 5/5 without needing the Responses API and the pricing page confirms the cost, or the decision log records why it stayed on `gpt-5.4-mini`.
DEPENDS ON	01

03 YOU 15 min

Hum into it with a real voice

The transcription was tuned on synthesized hums. The realistic generator is a careful model of a voice, but it is still a model.

HOW	Hum three takes: something slow with repeated notes, something fast, and a tune you know well. See Start here.
DONE WHEN	For each take you can say whether held notes shattered, whether repeated notes merged, and whether the tempo was right, half or double.
UNBLOCKS	04 and 05

HUM QUALITY

Making the transcription trustworthy

Everything here is measured against real takes, so it waits on step 03.

04 ○ HAND OFF small

Turn your hums into regression tests

Synthesized audio proves the pipeline works. Only labelled recordings of a voice prove it keeps working for one.

HOW	Save five to ten takes, each with the melody you meant written beside it, as fixtures next to the realistic suite. Keep them in an ignored folder if the repository will be public.
DONE WHEN	`uv run pytest` runs them alongside the existing 36 tests.
DEPENDS ON	03

05 ○ HAND OFF small to medium

Tune transcription on your takes

The loudness-dip prominence, the 80 ms minimum note and the 60–180 bpm tempo window were all set on synthesized audio.

DONE WHEN	Your recorded fixtures pass, and so do the synthetic tests.
DEPENDS ON	04, so every change is measured rather than judged by ear

06 ○ HAND OFF small

Fix tempo in one click

Half- and double-time readings are inherent to hearing a pulse. Correcting one today means editing the `setcpm` line by hand.

HOW	Add ×2 and ÷2 controls to the As hummed panel that rewrite `setcpm` and replay.
DONE WHEN	A take read at half tempo can be corrected without typing.

07 ○ HAND OFF small

Settle the tempo numbers and count the fallback

`_estimate_tempo`'s docstring says 19 of 21 realistic takes land within 10%, while the commit and `docs/humming.md` say 21 of 21. The `melody.py` docstring and the humming diagram still say tempo comes from `beat_track`. Nothing counts how often the `beat_track` fallback fires.

DONE WHEN	One number is true and quoted everywhere, the stale docstring and diagram are fixed, and the fallback path logs when it is used.

AGENT AND RELIABILITY

Fewer failures, lower latency

Independent of each other. Ordered by how soon they pay off.

08 ○ HAND OFF small

Retry model calls with backoff

A rate limit or quota error fails the whole request. The hum endpoint still returns its transcription, but the chat endpoint returns only an error.

DONE WHEN	A rate-limited call is retried with exponential backoff before anyone sees an error.

09 ○ HAND OFF small

Build the agent once

`run_llm` builds a new agent on every request, rebuilding its graph each time.

DONE WHEN	One agent is built at startup and reused, with latency measured before and after.

10 ○ HAND OFF medium

Keep a warm sandbox

Every validation starts Node and imports Strudel from scratch, and a correction loop pays that two or three times per answer.

TRADE-OFF	A long-lived worker gives up per-request isolation. Restart it every few dozen patterns and after any timeout, and keep the same permission flags.
DONE WHEN	Validation latency is measured before and after, and the file-write, network and subprocess escapes still fail inside the worker.

11 ○ HAND OFF small

Say so when the sandbox is missing

If `node` is not on the path or `sandbox/node_modules` is missing, the tool tells the agent to answer unverified and the server logs nothing.

DONE WHEN	The server logs a warning at startup, and the sandbox state is visible without reading an answer.

12 ○ HAND OFF small

Support the Responses API

`gpt-6-astra` and `gpt-5.6-sol` reject function tools on Chat Completions, and their replies arrive as content blocks that `run_llm` and the server treat as strings.

DONE WHEN	Both run through the full pipeline. Re-benchmark before switching, since neither beat `gpt-5.4-mini` on this task.

VALIDATION CORRECTNESS

Making VALID mean it plays

Three gaps between what the sandbox accepts and what the browser plays.

13 ○ HAND OFF small

Check sound names against the real sample list

The validator flags `gm_` soundfonts by prefix only. An invented drum name passes validation and then plays silence.

HOW	Fetch the dirt-samples sample map once, at install time, and ship it with the sandbox. The sandbox has no network by design, so it cannot fetch the list when it runs.
DONE WHEN	A pattern using `sound("bdd")` comes back with a warning naming the unknown sound.

14 ○ HAND OFF small to medium

Align the two Strudel versions

The page plays with `@strudel/web` 1.3.0, and the sandbox validates with `@strudel/core` 1.2.6. A function that exists in only one of them is judged by the other.

HOW	Find which core release the page's bundle contains. Then pin the page to a release that matches the sandbox, or run the sandbox on the page's own bundle if it will load in Node.
DONE WHEN	Both sides report the same version, and a test fails if they ever diverge.

15 ○ HAND OFF tiny

Revisit the `@kabelsalat/web` stub

The stub exists because every `@kabelsalat/web` release from 0.3.0 to 0.4.1 lacks `SalatRepl` and `@strudel/core` has nothing newer than 1.2.6.

DONE WHEN — Either a newer core or a fixed `@kabelsalat/web` is out and the stub is removed, or the check is repeated and the stub stays.

RETRIEVAL AND DATA

Keeping the corpus honest

BM25 and Pinecone are two views of one corpus, and they drift.

16 ○ HAND OFF small

Stop Pinecone drifting from the chunk dump

Ingestion overwrites `chunks.json` but upserts vectors, so a changed chunk leaves its old vector behind to keep surfacing in results.

DONE WHEN — Ingestion deletes vector ids that no longer appear in `chunks.json`.

17 ○ HAND OFF tiny

Decide whether `data/chunks.json` is committed

It is gitignored, but the README, `docs/extending.md` and the architecture traps say to commit it. A fresh clone has no dump, so retrieval quietly falls back to vector-only.

DONE WHEN — Either the file is committed and the ignore rule removed, or every doc says it must be generated first.

18 ○ HAND OFF small

Make a failed ingestion batch fail loudly

`add_batch` logs each batch's exception itself, so the closing success line prints and the process exits 0 even when batches failed. `chunks.json` is written first, so BM25 ends up ahead of Pinecone.

DONE WHEN — Any failed batch gives a non-zero exit and a summary of which batches failed.

19 ○ HAND OFF medium

Tune the hybrid weights on real queries

The 50/50 split and `k=10` per retriever are a reasonable default, not a measured one.

DONE WHEN — There is a small set of real questions and the weights are chosen against it. Needs usage data, so it waits on 33.

TESTS AND HOUSEKEEPING

Smaller things that add up

20 ○ HAND OFF medium

Add tests beyond `melody.py`

The agent, retrieval, server and frontend have no tests, so checking them means a manual run with API keys.

HOW — Start with the server endpoints using a stubbed `run_llm`, then `repair_answer` and the sandbox report parser.

DONE WHEN — `uv run pytest` exercises the server without network or keys.

21 ○ HAND OFF tiny

Fix `.env.example`

It describes an `OPENAI_CHAT_MODEL` variable and a `search_web` tool that do not exist, and marks `PINECONE_API_KEY` optional when the app needs it.

DONE WHEN — It matches what the code reads.

22 ○ HAND OFF tiny

Fix the ingestion docstring

`ingestion.py`'s module docstring says it crawls `python.langchain.com`. It crawls `strudel.cc`.

DONE WHEN — The docstring is correct.

23 ● YOU 15 min

Confirm or delete the inferred decisions

Several entries in `docs/decisions.md` are marked "inferred — confirm": hybrid retrieval weights, FastAPI over Flask, Pinecone over Chroma, and the vanilla frontend over Gradio.

DONE WHEN — Each is confirmed with its real reason, or deleted.

24 ● YOU about 1 h

Review and merge the branch

Thirteen commits of new subsystems sit on `feat/hum-to-pattern` and `main` has none of them. Each message explains why, so reading them in order is the fastest way in.

HOW — `gh pr create --base main --head feat/hum-to-pattern`. Merging after the reliability steps keeps the diff reviewable.

DONE WHEN — The branch is merged into `main`.

LATER

Before anyone else uses it

None of these matter while the app runs on your own machine. All of them matter the day it does not.

25 ○ HAND OFF medium

Pick a host and serve over HTTPS

`python app/server.py` binds localhost only and nothing in the repo says where it would run. `docs/runbook.md` has a TODO where the deploy steps should be.

DONE WHEN — The backend has a real hostname and certificate, and the runbook has real deploy, rollback and what-to-check-when-it's-down steps.

26 ● YOU a conversation with a lawyer

Get the AGPL question answered

Strudel is AGPL-3.0. The sandbox runs it inside a networked service, which is the case the licence's network clause addresses.

DONE WHEN — You know what source you must publish.

27 ○ HAND OFF small

Add rate limits and a spend cap

Both endpoints call paid APIs and neither is limited, so an open instance is an open tap on the OpenAI bill.

DONE WHEN — Requests are limited per client and a daily spend ceiling stops the model calls.

IDEAS

Worth exploring, not yet worth planning

28 · Lyrics from audio

Record or attach audio and ask about the lyrics. Needs a speech-to-text step, because today the

29 · Spoken instructions

Audio models failed at pitch, but transcribing words is what they are good at. Saying "make it

agent never hears the audio and `/api/chat` takes text only. The text would feed the chat flow beside the melody summary.

darker" is a natural gesture. Shares the speech-to-text step with 28.

30 · Chords and harmony

The pipeline is monophonic. A polyphonic transcriber such as Spotify's basic-pitch could add harmony, if its dependencies install on Python 3.13.

31 · Swing and triplets

The grid is straight sixteenth notes, so a shuffle is forced onto the nearest step.

32 · A native app again

The Capacitor shell was removed on 14 October to focus on the web app. It is recoverable from git history (`2d88161`) and only makes sense once the backend has a deployment target.

33 · Measure it in use

LangSmith tracing is already wired through environment variables, but nothing records latency, tool calls per answer, or how often validation changes the model's pattern.

Grounded in the code at commit `9814eb9`. For how the pieces fit together, see docs/architecture.md.